



Running LLMs on Kubernetes?

Cut cost without sacrificing GPU performance and response time

Le

MATHEMATICS

Desmystifying Curvelets

Learn what curvelets are what they can be used for. I about signal processing and computing computing.

Carlos Costa, Ph.D.

Mar 22, 2023 16 min read

LATEST

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

Signal-processing deep dive

Learn what curvelets are, how they are built and be used for

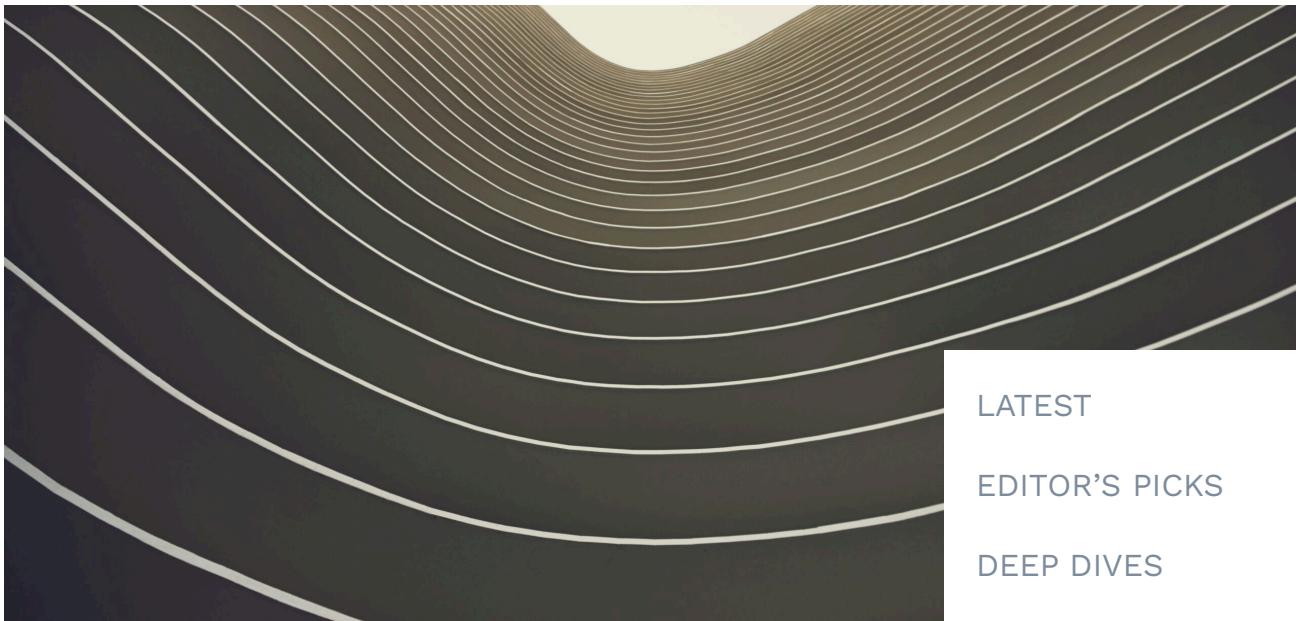
What AI Roles Are Organizations Hiring?

https://towardsdatascience.com/desmystifying-curvelets-c6d88faba0bf/?source=social.linkedin&_nonce=FkXhiRkE

1/26



What AI roles are growing? 57% are hiring AI engineers, +11% increase from 2024. Learn more about the value of human talent.



Edifício Niemeyer, Belo Horizonte. Photo by [Matheus Frade](#) or

LATEST

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

WRITE FOR TDS

[Sign in](#)

[Submit an Article](#)

Introduction

Curvelets are multiscale, oriented, non-adaptive of images and multi-dimensional signals. If something didn't make sense, you're in the right place.

Developed in the early 2000s by Emmanuel Candès and Donoho [1] in the flurry of wavelet-transform related processing boom, curvelets were designed to solve the issues that plagued previous alternatives. Wavelets generalize the Fourier transform by using arbitrary – but especially smooth – functions instead of complex exponentials. Like the Fourier transform, they can be readily extended to 2D by repeatedly applying the transform over multiple scales. Constructing 2D wavelets with this naïve approach leads to issues when representing edges that are not exactly vertical. In practice, transforms that have

edges can cause "blocky" artifacts in strongly compressed images. Many transforms suffer from the same fate, including the discrete cosine transform (DCT) which powers the ubiquitous JPEG format (see Figure 1). JPEG2000 which relies on the wavelet transform to improve on JPEG, still suffers from the same blocky artifacts.



Figure 1. "Blocky" artifacts appearing in strongly compressed JPEG images for b
Credit to: [Shlomi Tal \(CC BY-SA 3.0\)](#).

LATEST
EDITOR'S PICKS
DEEP DIVES
NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

Curvelets don't have these issues – *they can pre* edges. They also exhibit some nice properties, s optimal representation of wave-like phenomena conservation, unitarity (having its adjoint/transp as its inverse), among others.

Today, curvelets are used in a variety of image-p including denoising, compression, inpainting, sm the geophysical community, it has become a po' for several tasks including adaptive subtraction, preconditioning, among others. In the medical c been used for segmentation, diagnosis, etc.

While it is true that deep learning has quickly di algorithms for many of these tasks, curvelets (and transforms) still have their use. Whereas deep le

possibility to create *adaptive*, multiscale representations of signals, curvelets already offer that predictably, without requiring training, and optimally for many types of signals. One may use curvelets directly as a substitute when no training data is available, or even employing them alongside deep learning methods, thereby reducing model complexity and data requirements.

In this deep dive, we will go over the building blocks of the curvelet transform, focusing on the intuition behind it, and how it can be used to provide a starting point for users interested in applications which use curvelets. Find all the code and more in the [Curvelops repository](#).

[LATEST](#)[EDITOR'S PICKS](#)[DEEP DIVES](#)[NEWSLETTER](#)[WRITE FOR TDS](#)[Sign in](#)[Submit an Article](#)

Before Curvelets: the Fourier Transform

To understand the curvelet transform, we need to understand the Fourier transform. The 2D FFT (Fast Fourier Transform) tells us what kind of spatial frequencies are present in our image. The more energy on the larger (away from the origin) wavenumbers (spatial frequencies), the faster the signal varies along a certain direction. In this section we will learn how to find this direction, as well as how to quantify it.

Let's start by creating images which vary rapidly in one particular direction, but not at all in the perpendicular direction. These will be helpful in understanding what the spectrum (another name for Fourier domain, the domain after applying FFT) is telling us. We will create an image modulating a cosine in the direction normal to a line ($\cos \theta$).

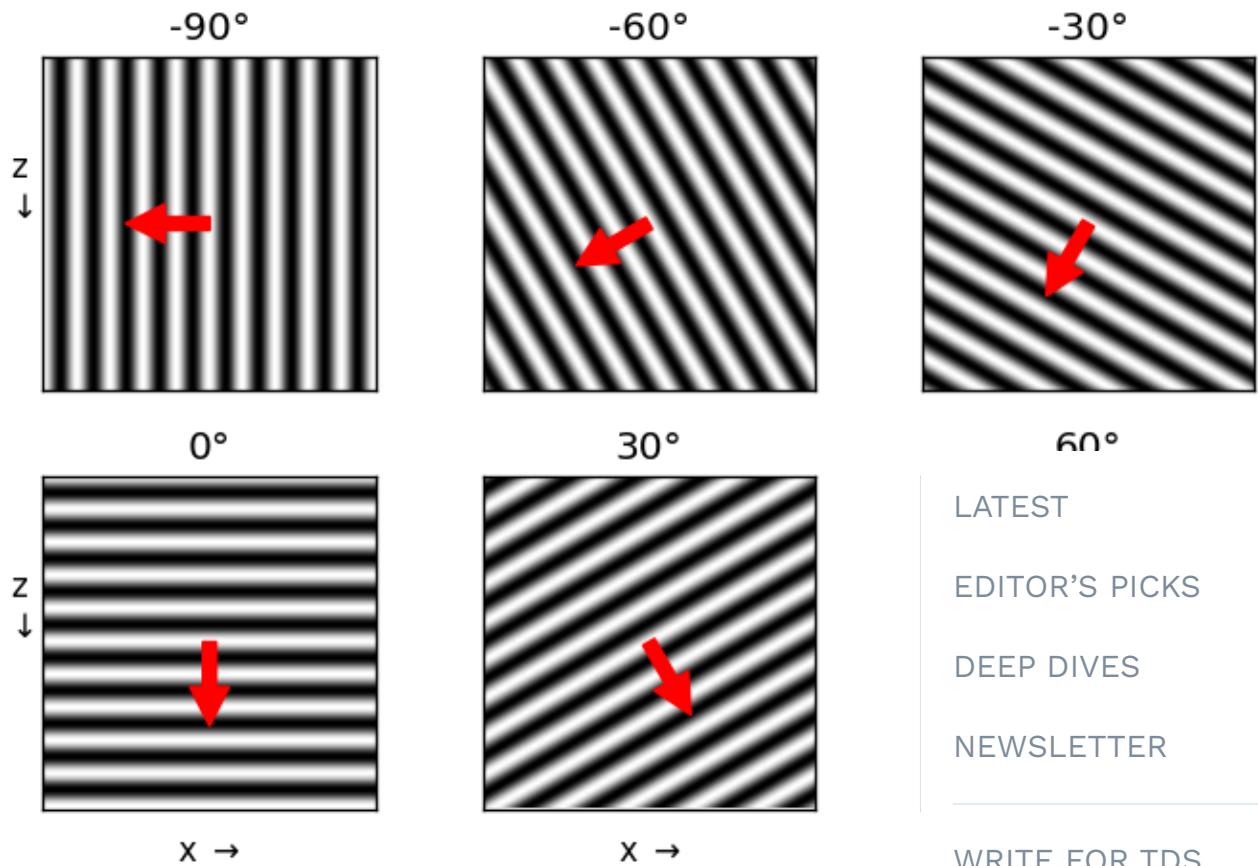


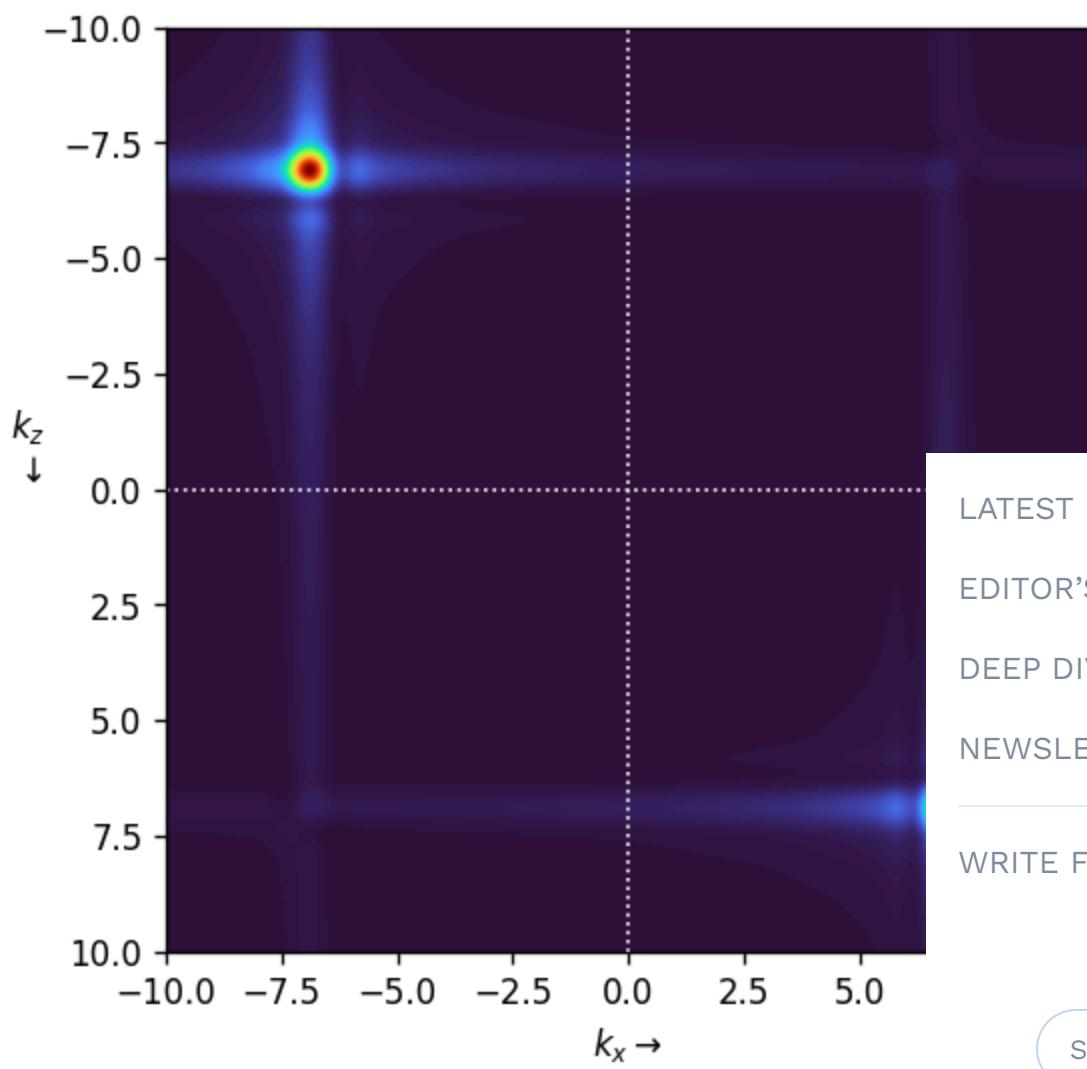
Figure 2. 2D monochromatic images, with direction of change in each image derived from work.

[Sign in](#)

[Submit an Article](#)

Each plot in this image has an arrow pointing to the direction of maximum variation v . Along the direction perpendicular to it, the signal does not vary at all.

This won't always happen, signals can vary in more than one direction at once. The question is, which directions? We will find out? Cue the FFT transform. Let's take our 2D images, say the one at 45°, and apply the 2D FFT.



LATEST
EDITOR'S PICKS
DEEP DIVES
NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

Figure 3. 2D FFT of the monochromatic signal which varies most at the 45° c

In the Fourier domain, we can see that our signal is symmetric about the origin. A property of the 2D Fourier transform for real signals is that it is Hermitian, meaning that its Fourier spectrum is symmetric about the origin. In the above plot, the top-left quadrant carries the same information as the bottom-right quadrant; the top-right carries the same information as the bottom-left quadrant. So we can "ignore" the negative frequencies for one axis. One way to choose the "fastest" axis, in this case the vertical axis.

WARNING: This is not true for complex signals. It only changes the math much, just the visual interpretation.

only use real signals here.

So we can scan the (positive k_z) k -space to identify a single peak. Its corresponding k_x and k_z coordinates are given by: [6.931, 6.931]. Not particularly interesting, until we normalize this vector to [0.707, 0.707] and compare it with the original direction of maximum variation: [0.707, 0.707]. Up to a constant, the vectors are the same! Let's try this for all images create above

LATEST

1. Compute the 2D FFT
2. Find the maximum amplitude location in k -space
3. $(k_x\text{-max}, k_z\text{-max})$ is the direction of maximum variation in the input

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

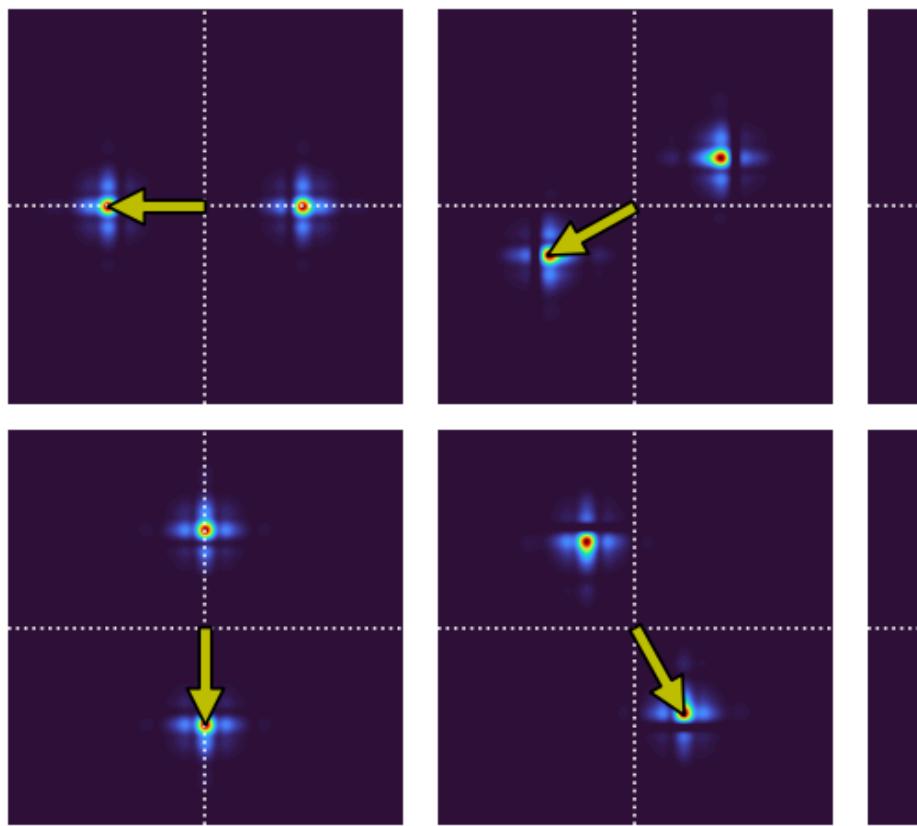


Figure 4. 2D FFT of the monochromatic signals in Figure 2. Create

Notice how these follow exactly the direction of maximum change in the data. Indeed, if we overlay these normalized vector onto the data-space images, we get:

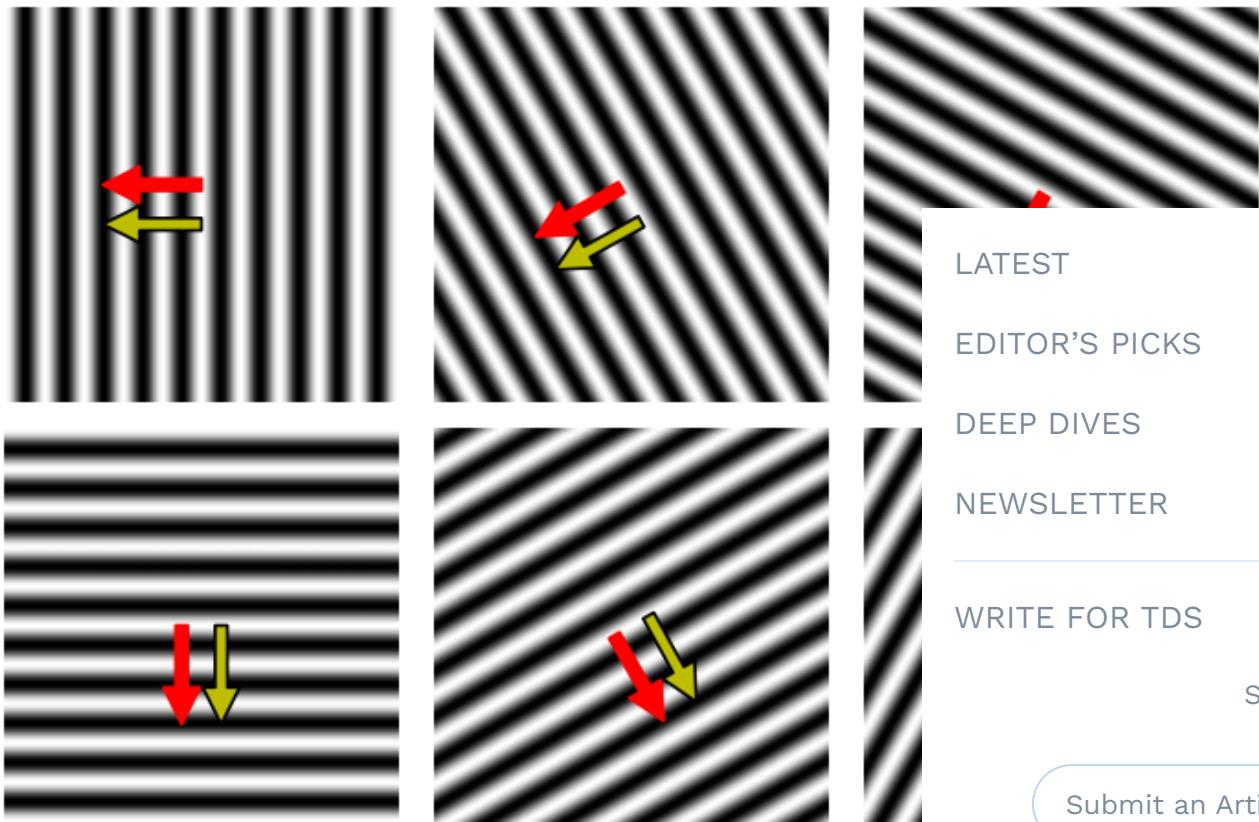
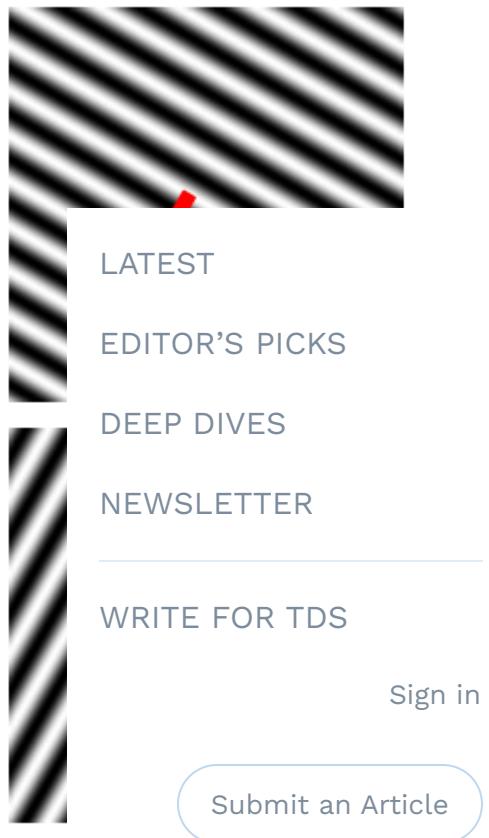


Figure 5. Monochromatic signals in Figure 2. Red arrow: direction of maximum variability estimated from 2D FFT. Credit: Author

Looks good! What does this have to do with the transform? Well, imagine we try applying the same following signal:



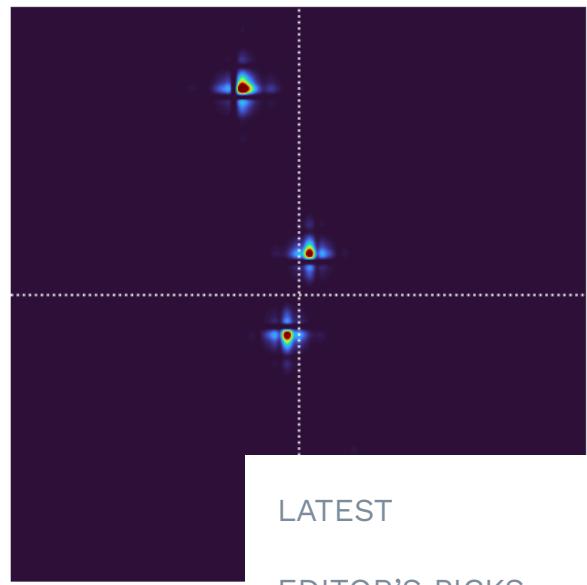
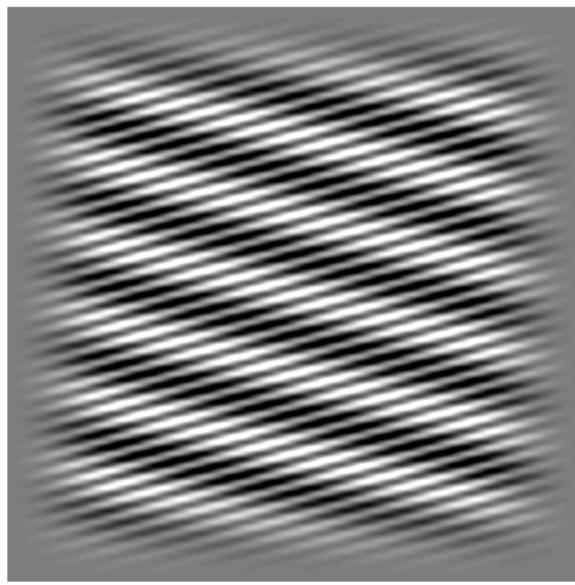


Figure 5. Bichromatic signal in spatial (left) and Fourier (right) domain

We can see that there is a low-frequency signal background, corresponding to the -15° component. There is a high frequency component in another direction. If we apply our k -max algorithm, because the low-frequency component is weaker, it won't get picked up. We will think that the dominant direction is 15° . But the FFT spectrum has all three components. So what can we do with it?

LATEST

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

Building the Curvelet Transform

Separating scales

The curvelet transform handles this wealth of information by separating the signal at different *scales*, which are concentric regions in the k -space domain. These regions encompass frequencies that are similar. Let's apply this to the example above:

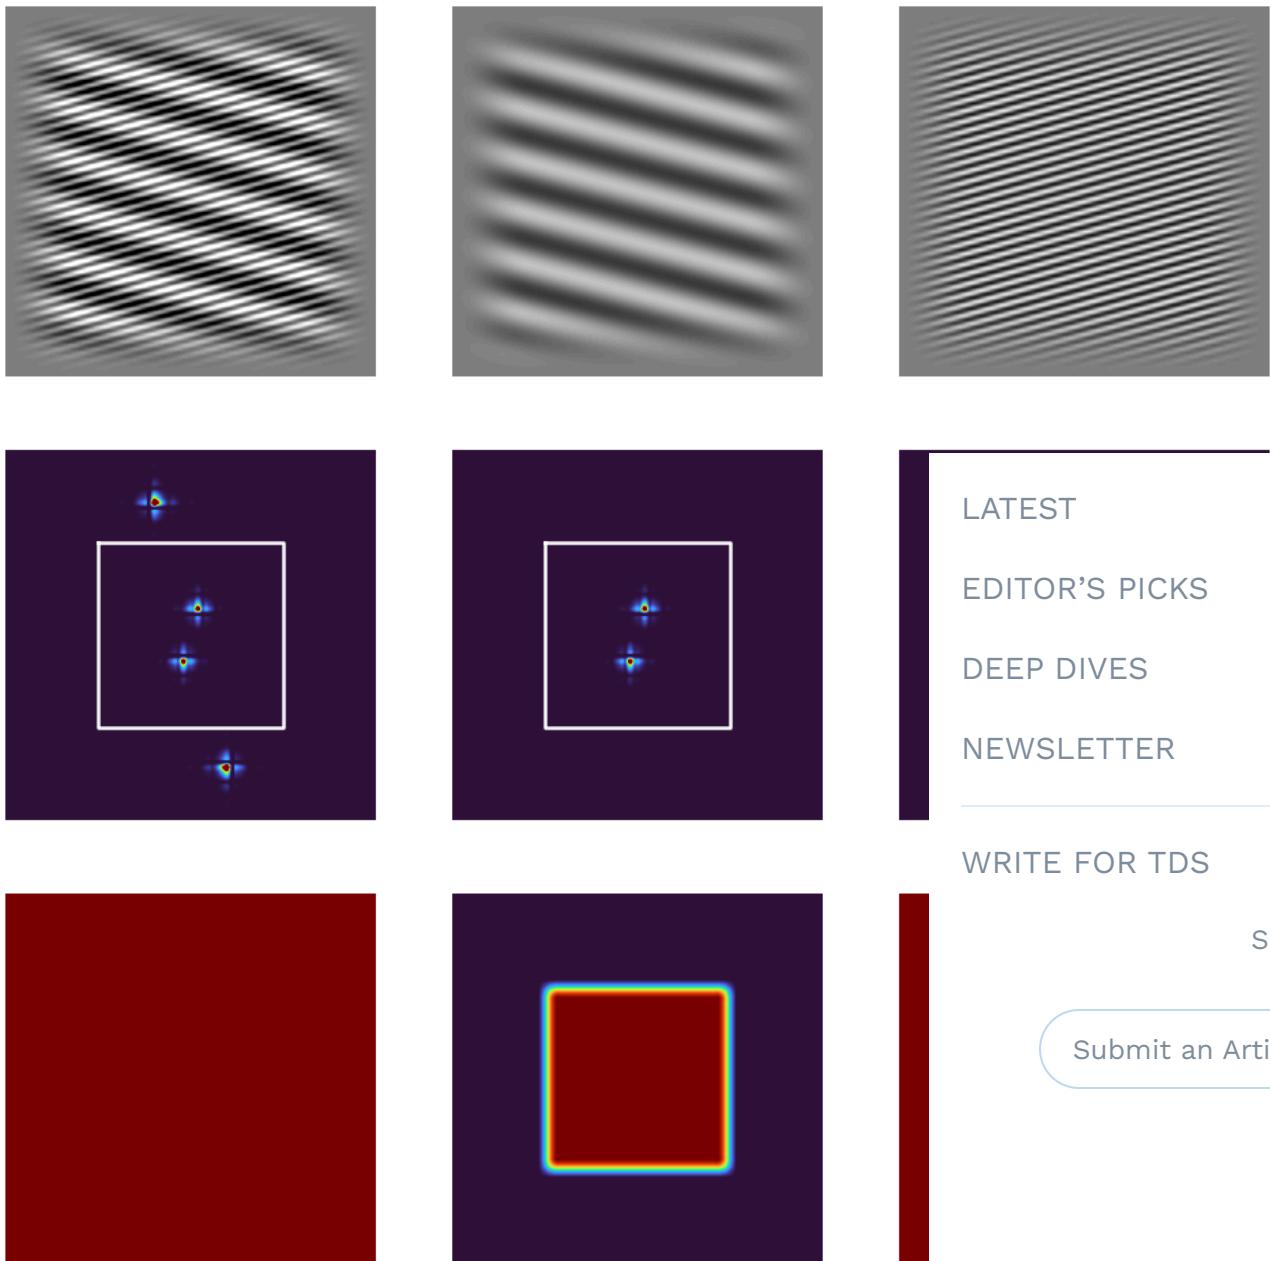


Figure 6. Original bichromatic signal (left column). Low-passed signal (center column). Reconstructed signal (right column). Spatial domain (top row). Fourier domain (middle row). Mask generate the signals in the top row (bottom row). Deep purple denotes zero values, red denotes unity. Credit: own work.

The top-left signal is the original input signal. The top-center (second row) is its the Fourier spectrum. The image in the bottom row (third row) is the "mask" applied to the Fourier spectrum to reconstruct the original signal. It is entirely red, ones, that is, no change to the spectrum.

[LATEST](#)
[EDITOR'S PICKS](#)
[DEEP DIVES](#)
[NEWSLETTER](#)

[WRITE FOR TDS](#)
[Sign in](#)

[Submit an Article](#)

The second column, depicts signals from the first region, constructed by a low-pass filter defined by its mask (third row, middle column). Where it zeroes the signal, the mask is purple. Finally, the third column depicts the second region, constructed by a high-pass filter which is $1 - \text{lowpass}$.

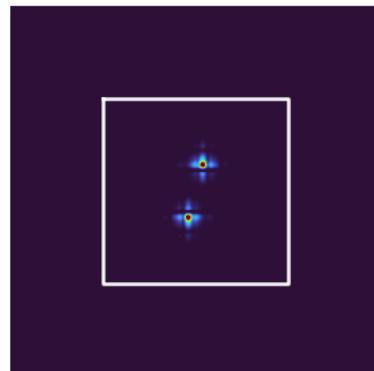
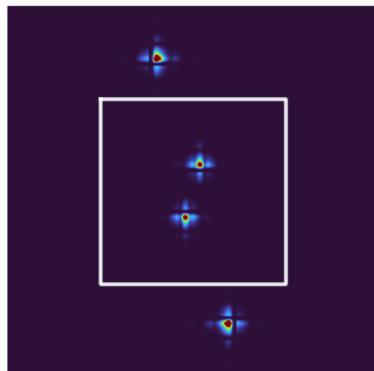
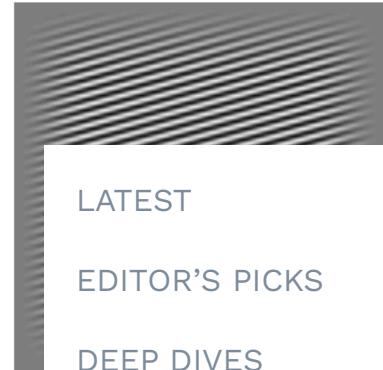
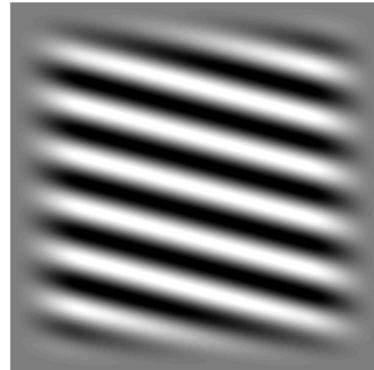
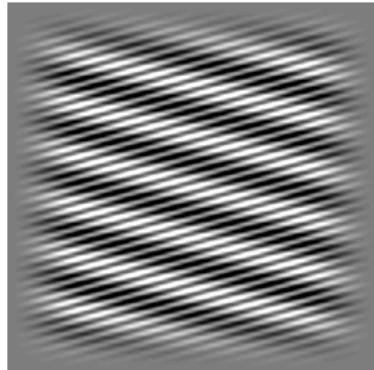
In the curvelet transform, the number of scales is the first parameter. A choice of scales = 2 would create a division very similar to the one above (but with a different, slight smoothing). The more scales there are, the more curvelet transform coefficients will be able to separate with different frequencies.

And what will these coefficients be? In this case similar to the inverse transform of the low-pass filters, that is, the two rightmost panels in the top row. A caveat (in addition to the smoothing) is that the scale does not require the same sampling as the original. Indeed, the Nyquist frequency (highest frequency sampling can represent) of the lowpass signal is half the original Nyquist frequency. This is clear from the fact that the lowpass signal removes all frequencies above Nyquist. Therefore we could actually *double* the sampling component compared to that of the original signal. As a general rule, and we will see that the coarsest scale can be sampled at 2^m times the original sampling rate, where m is equal to the number of scales minus 1.

The finest scale (in this example, scale = 1) can be sampled in the same way, as its highest frequency remains the same as the original signal.

[LATEST](#)[EDITOR'S PICKS](#)[DEEP DIVES](#)[NEWSLETTER](#)[WRITE FOR TDS](#)[Sign in](#)[Submit an Article](#)

For visualization's sake, we can apply this resampling to the lower frequency scale to get an idea of what the curvelet coefficients would look like for a scales = 2 and no additional subdivisions (more about this soon).



LATEST
EDITOR'S PICKS
DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

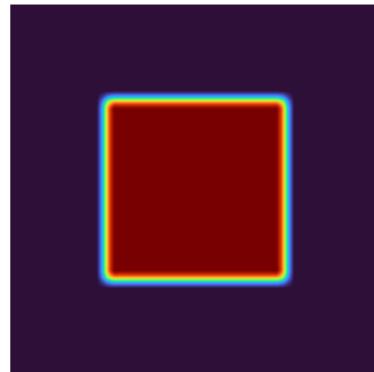
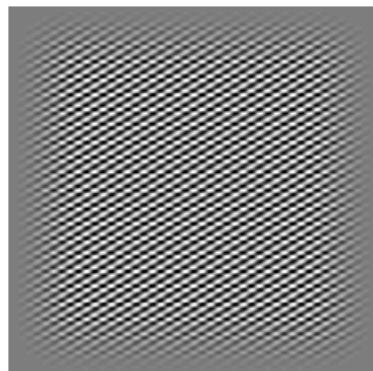


Figure 7. Original signal (left column). Coarsest curvelet scale using scales = curvelet scale (right column). All other panels like Figure 6. Note: this curvelet t finest scale. Credit: own work.

Comparing this plot to the original plot without notice that the signal is accurately represented

strength of the resampled signal is larger. This is because the Fourier transform used to "return" to the spatial domain now has a normalization factor which is smaller than the original. Therefore the signal will be twice as strong ($2 = \sqrt{2} \times \sqrt{2}$, one square root per dimension, in this case two). This is also a general rule of the curvelet transform and Fourier transforms.

Let's apply this decomposition to another example:



LATEST

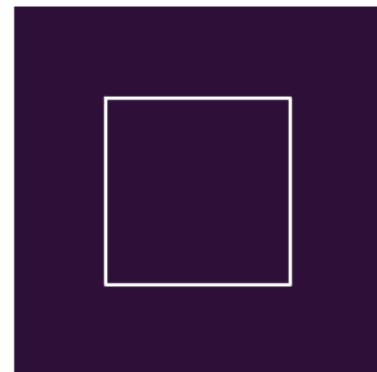
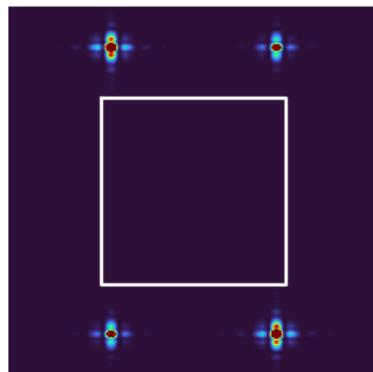
EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in



Submit an Article

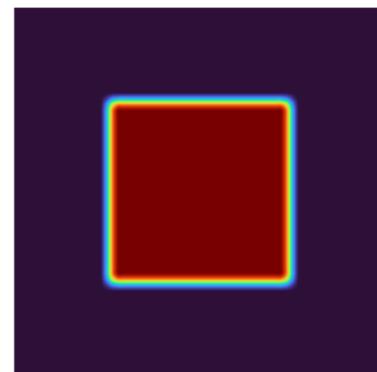


Figure 8. Tetrachromatic signal composed of four signals: two spatial freq orientations. Panels like Figure 7. Credit: own work

Separating Dips

While we have separated in scale, we still are not able to distinguish the two separate events, one at -30° and one at 30° . We need another type subdivision, one which splits a signal (of similar frequencies) further, into "subsignals" of similar *direction* (or *dip*, in geophysical lingo). Starting with the second scale (the one after the coarsest), we will further split it ~~scale into wedges~~ These are approximately angular sectors in the i

LATEST

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

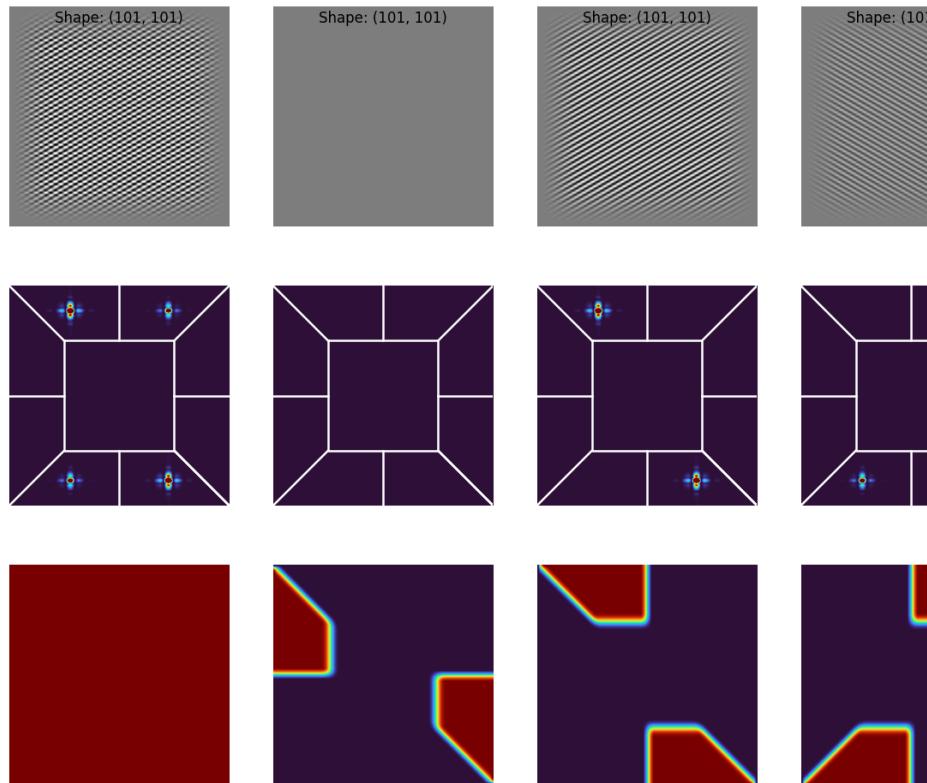


Figure 9. Original tetrachromatic signal (leftmost column). Signal separated by "center-left to rightmost columns". Rows like Figures 6, 7 and 8. Ci

We can see that we have been able to separate two constituent components. For this signal, two are identically zero. The signals in the first row, second column, are essentially what the curvelets

would look like if we applied the curvelet transform with the parameters scales = 2 and wedges = 8.

One small caveat is that we are actually only showing **four wedges**. This is again because of the symmetry of the Fourier domain for real signals. For complex signals, we would have to treat each wedge (without its symmetrical counterpart) separately.

LATEST

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

We now have all ingredients to understand how Transforms are implemented in practice. For example, we have some kinks to iron out. One of them is the fact that when we sum every wedge mask, we don't get an exact value at every grid point in the k -space domain. This means that we either "lose" or "create" signal. A more pressing issue is that the wedges were transformed back to the spatial domain in a non-optimal way: we used a full-sized FFT transform to transform a small sliver of the k -space domain. All curvelet coefficients all shaped the same shape and size as the input image (in that example, 101×101).

Moreover, an important property of the continuous curvelet transform is that performing the steps backwards, i.e., outputting an image *from* curvelet coefficients, should yield the exact same image that was used to obtain these coefficients in the first place. Mathematically, the transform is unitary, which means that its adjoint is equal to its inverse. Our implementation respects this property.

All these issues can be fixed. One construction which solves all these problems is the Fast Discrete Curvelet Transforms, developed by Emmanuel Candès, Laurent Demanet, David Donoho and Lexing Ying. Its most famous implementation is provided by the [CurveLab package](#), which the [Curvelops](#) package wraps via Python (disclaimer: I developed Curvelops). Curvelops thereby provides a linear operator interface relying on [PyLops](#) (disclaimer: I am one of the core devs of this lib

LATEST

Let's explore the FDCT by applying it to the sign compare the curvelet coefficients with what we before.

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

Scale 1 (4 wedges)

Wedge 0

Shape: (135, 59)



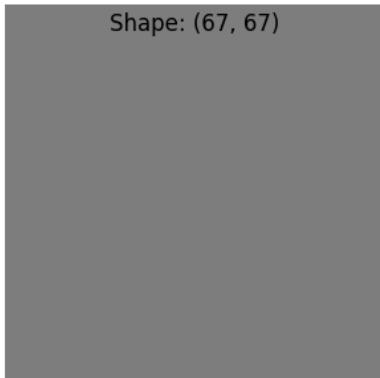
WRITE FOR TDS

Sign in

Submit an Article

Scale 0 (1 wedge)

Shape: (67, 67)



Wedge 2

Shape: (59, 135)

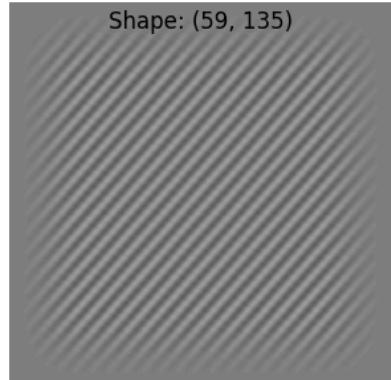


Figure 10. (Real) curvelet coefficients of the tetrachromatic signal. ↗

These look very similar to what we obtained before, with a few differences. First, as mentioned before, the algorithm actually outputted 8 wedges in the second scale; we merged them due to the symmetry of the Fourier space for real input signals. Second, we can see that the shapes are not the same as the input signal. The coarsest scale is about $\frac{2}{3}$ of the input signal. On the second scale, not only the shapes are different, but the aspect ratio is also different. This is because different directions have different sampling. The FDCT handles all these in a way that is performant and that respects the underlying continuous transform.

[LATEST](#)[EDITOR'S PICKS](#)[DEEP DIVES](#)[NEWSLETTER](#)[WRITE FOR TDS](#)[Sign in](#)[Submit an Article](#)

Now we are ready to understand the parameters:

- **nbscales** (no default, minimum 2): Number of scales. The more scales, the finer the details the curvelets capture. The downsides to more scales are tapering issues. Both of these issues can be avoided by setting `allcurvelets=False`.
- **nbangles_coarse** (default: 16, minimum 8, maximum of 4): Number of angles/wedges at the second scale. Remember that the coarsest scale is never included. This value is only specified for the second scale because it is doubled every two scales (counting as if there were 4 nbangles_coarse wedges). So by setting a `nbscales=4` and `nbangles_coarse=8`, the second scale will have 16, the third 16 (as we have seen 2 scales by now), the fourth and last, 32 (as we have seen four scales).

Downsides to increasing this parameter is a

and tapering issues. Mitigate them also with `allcurvelets=False`.

- **allcurvelets** (default: True, boolean): Controls whether we perform wedge subdivision on the finest (last) scale. We did this in the latter example, but not in the one before that. From a practical point of view, despite defaulting to True, you should set this options to False unless there is a strong reason not to.

LATEST

Examples of the Fast Discrete Curvelet Transform

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article



Figure 11. Python "two-snakes" Logo. Credit: [Python Software](#)

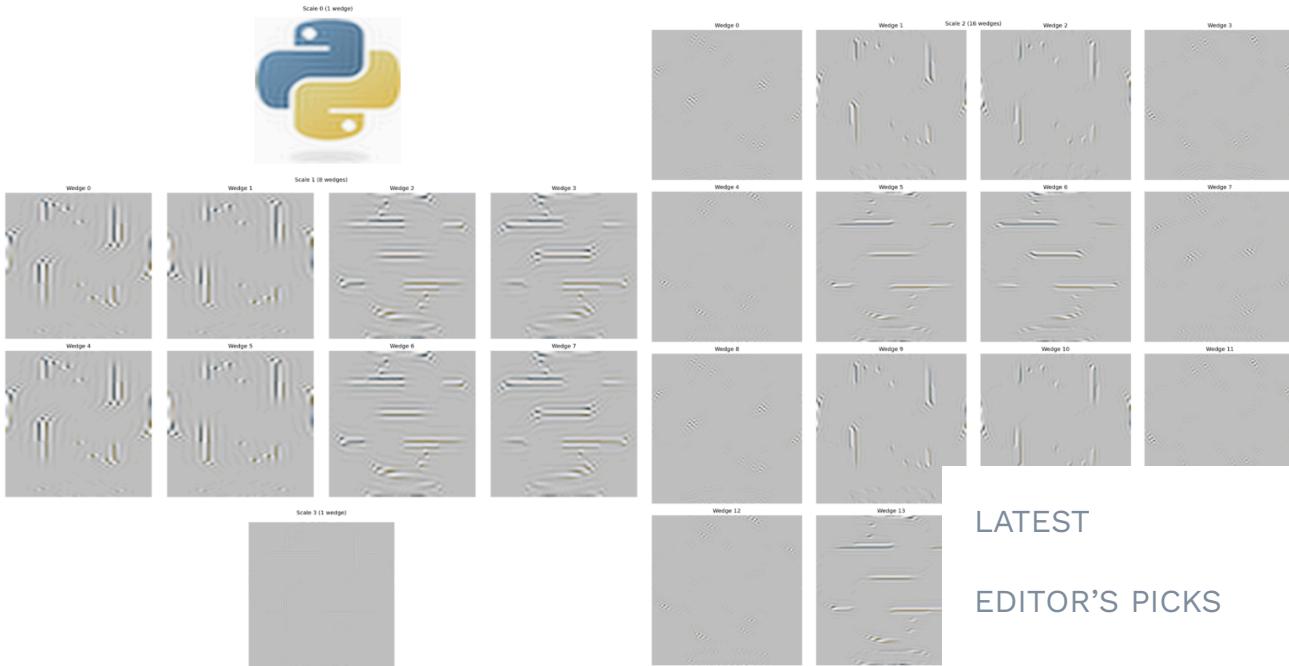


Figure 12. Curvelet coefficients (nbscales=4, nbangles_coarse=8, allcurvelets=False)
Figure 11. Credit: own work.

LATEST
EDITOR'S PICKS
DEEP DIVES
NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

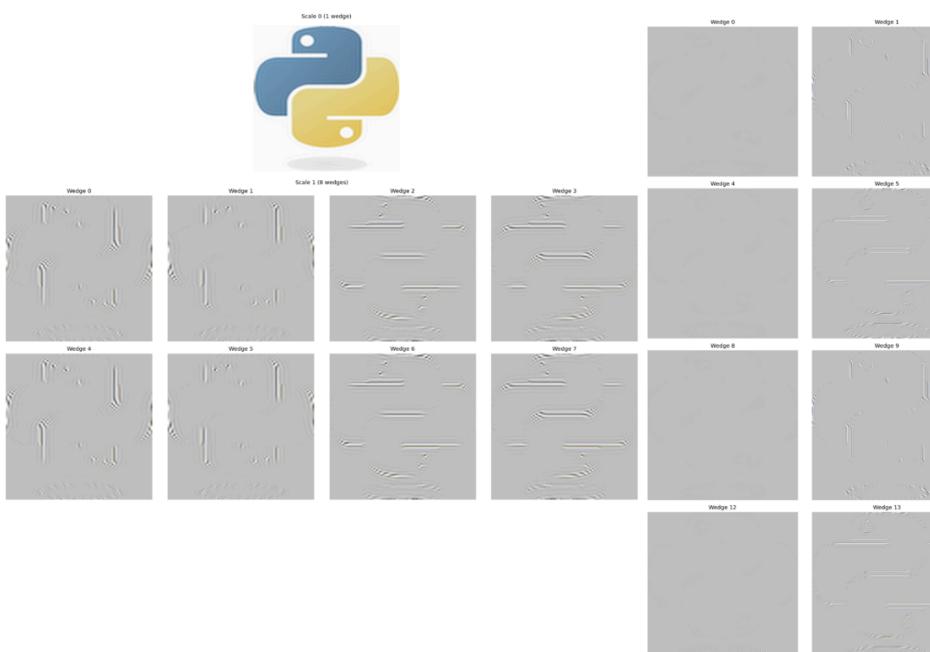


Figure 13. Curvelet coefficients (nbscales=3, nbangles_coarse=8, allcurvelets=True)
Figure 11. Credit: own work.

It is interesting to notice that the "color" of the the coarsest scale because this scale contains the frequency aka the DC-component. All other scale variations on top of that initial image.

Another interesting aspect of the curvelet transform is that it is very easy to interpret its coefficients. They are essentially a "piece" of the original image which varies along certain preferential directions.

But visualizing all of these coefficients can quickly become overwhelming. One way to overcome that is to extract features from a single wedge. In the following example we will subdivide each wedge into rows/cols, of which we will consider LATEST These energies will be mapped to disks in the EDITOR'S PICKS will see below how this visualization can help us DEEP DIVES preferential directions of events in an image in the NEWSLETTER

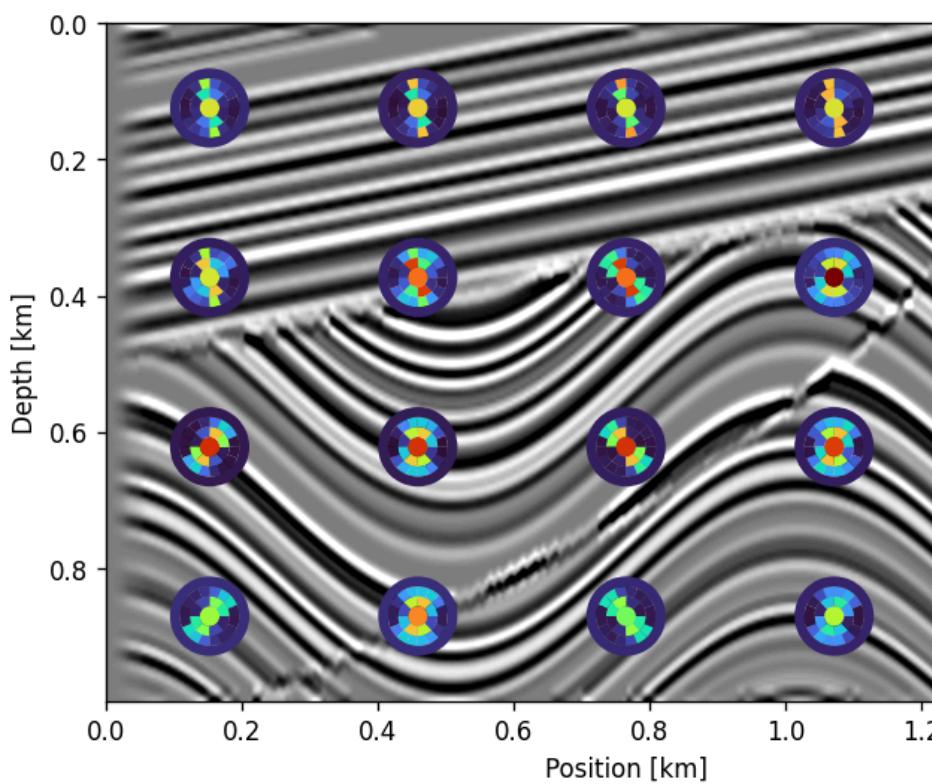


Figure 14. Sigmoid seismic model with curvelet strength disks overlain. Strength is perpendicular to the local dip. Credit: own work.

Here we are plotting, overlain on the original input seismic model) the strength of the curvelet coefficients in a certain window in each scale/wedge. These disk

wedge division we discussed above: the closest to the center, the lower the scale; the angular wedges map to the same place in the k -space domain, in a polar projection.

Consequently, similar to our visualization of the k -max vectors, we should see that the strongest energy lies *perpendicular* to the preferential local dips. As a reminder, this happens because the strongest points in the k -space are those where ~~the image varies~~ the most in that direction. So, perpendicularly to ~~the image varies~~ the least. We can identify this behavior in the image near the top of the image, where the structure is ~~the image varies~~.

LATEST

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

However, these disks give us even more information about the preferential dip at any arbitrary location. They also give us an idea of the anisotropy of the image. For example, if there is no preferential direction, or many preferential directions, the image is more isotropic. And this information is not only locally in space, but also separated by scales (smaller vs larger frequencies).

Of course, even the disks are a aggregate of the smaller components that the curvelet transform contains, which attest to its power and versatility!

These are just some examples of the curvelet transform's power. In upcoming articles, we will explore how they can be used for various tasks!

Key Takeaways

- The **Fourier transform** of images give us an idea of the **preferential directions of change**
- The **curvelet transform** goes a step beyond that, telling us how the image is **varies at each location, in which direction and with which spatial frequency**.
- The curvelet transform is traditionally performed with the FDCT (**Fast Discrete Curvelet Transform**) provided in **Python by curveloops and CurveLab**LATEST
- The FDCT can be used in many areas such as **image processing and deep learning** (see TorchOpener) EDITOR'S PICKS
- DEEP DIVES

NEWSLETTER

WRITE FOR TDS

Sign in

Submit an Article

Further Reading

[1] Candès, E., Demanet, L., Donoho, D., & Ying, L. (2005). Discrete Curvelet Transforms. *Multiscale Modeling and Simulation*, 5(3), 861–899.

[2] Ma, J., & Plonka, G. (2010). The Curvelet Transform. *Image Processing Magazine*, 27(2), 118–133.

• • •

WRITTEN BY

Carlos Costa, Ph.D.

See all from Carlos Costa, Ph.D.

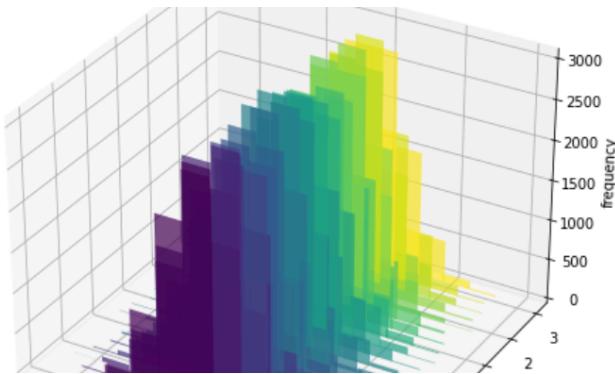
[Image Processing](#)[Mathematics](#)[Scientific Computing](#)[Signal Processing](#)[Wavelet](#)**Share This Article**

Towards Data Science is a community publication where authors share their insights to reach our global audience and receive payment through the TDS Author Payment Program.

[LATEST](#)[EDITOR'S PICKS](#)[DEEP DIVES](#)[NEWSLETTER](#)[WRITE FOR TDS](#)[Sign in](#)

Related Articles

[Submit an Article](#)



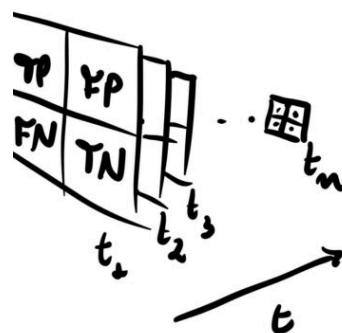
DATA SCIENCE

Must-Know in Statistics: The Bivariate Normal Projection Explained

Derivation and practical examples of this powerful concept

Luigi Battistoni

August 14, 2024 7 min read



$$= p(\hat{y}, y, t)$$

MACHINE LEARNING

Binary Classification

Unpacking the Results of a Machine Learning Model

Limitations of Traditional ML Models

gabriel costa

January 12, 2024

LATEST

EDITOR'S PICKS

DEEP DIVES

NEWSLETTER



DATA SCIENCE

Let's Talk About Math (for Data Scientists)

Our weekly selection of must-read Editor's Picks and original features

TDS Editors

March 23, 2023 3 min read

ARTIFICIAL INTELLIGENCE

Teach (and learn) by transforming geometry

Understanding the complex plane through a visualization

Peter Barrett Bryan

March 18, 2022 3 min read

Sign in

Submit an Article



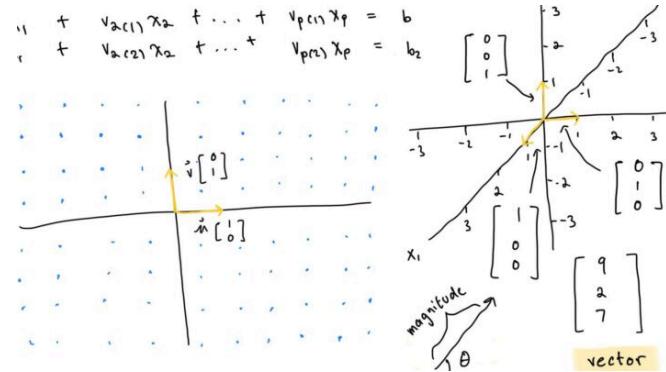
DATA SCIENCE

Multi-Dimensional Exploration Is Possible!

(At least mathematically)

Diego Manfre

October 5, 2023 16 min read



MACHINE LEARNING

Linear Algebra Equations

LATEST

Vector Equations

EDITOR'S PICKS

Tenzin Migmar

October 14, 2023

DEEP DIVES

MACHINE LEARNING

COUNTLESS 3D- Vectorized 2x Downsampling of Labeled Volume Images Using Python and Numpy

Previously, I demonstrated a fully vectorized algorithm, COUNTLESS, that downsampled labeled images by finding the...

William Silversmith

February 20, 2018 19 min read

WRITE FOR TDS

Sign in

Submit an Article



Your home for data science and AI. The world's leading publication in analytics, data engineering, machine learning, and artificial intelligence.

[Subscribe to Our Newsletter](#)[WRITE FOR TDS](#) · [ABOUT](#) · [ADVERTISE](#) · [PRIVACY POLICY](#) · [TERMS OF USE](#)[COOKIES SETTINGS](#)[LATEST](#)[EDITOR'S PICKS](#)[DEEP DIVES](#)[NEWSLETTER](#)[WRITE FOR TDS](#)[Sign in](#)[Submit an Article](#)